

데이터베이스 #assignment5

수업		Cose371-02		교수		정연돈교수	
대학	정보대학	학과	컴퓨터과	학번	2018320211	이름	진지과

Index Creation and Scan

Q1.create index

Q2.Make queries eachuses 'seq scan', 'index scan', and 'index only scan'

seq scan

index scan

index only

Q3.

Q4.

B-tree and Hash

create two indexes

Q5.

Q6.

Q7.

Q8

Q9.

Index Creation and Scan

Q1.create index

```
create index index on table1 (sorted); //clustered index
create index index on table1 (unsorted); //non-clustered index
```

Q2.Make queries eachuses 'seq scan', 'index scan', and 'index only scan'

seq scan

```
explain analyze select * from table1;
```

```
[hw5=# explain analyze select * from table1;
                                QUERY PLAN
-----
Seq Scan on table1 (cost=0.00..203093.00 rows=10000000 width=53) (actual time=1.829..1406.920 rows=10000000 loops=1)
Planning Time: 0.418 ms
Execution Time: 1776.791 ms
(3 rows)]
```

index scan

```
explain analyze select rndm, sorted from table1 where sorted=320211;
explain analyze select unrndm, sorted from table1 where unsorted=320211;
```

```
hw5=# explain analyze select rndm, sorted from table1 where sorted=320211;
                                QUERY PLAN
-----
Index Scan using table1_sorted_idx on table1 (cost=0.43..8.70 rows=15 width=8) (actual time=0.023..0.027 rows=5 loops=1)
  Index Cond: (sorted = 320211)
  Planning Time: 0.114 ms
  Execution Time: 0.049 ms
(4 rows)

hw5=# explain analyze select rndm, unsorted from table1 where unsorted=320211;
                                QUERY PLAN
-----
Index Scan using table1_unsorted_idx on table1 (cost=0.43..28.54 rows=6 width=8) (actual time=0.051..0.055 rows=4 loops=1)
  Index Cond: (unsorted = 320211)
  Planning Time: 0.065 ms
  Execution Time: 0.068 ms
(4 rows)
```

index only

```
explain analyze select sorted from table1 where sorted=320211;
explain analyze select unsorted from table1 where unsorted=320211;
```

```
hw5=# explain analyze select sorted from table1 where sorted=320211;
                                QUERY PLAN
-----
Index Only Scan using table1_sorted_idx on table1 (cost=0.43..8.70 rows=15 width=4) (actual time=1.307..1.312 rows=5 loops=1)
  Index Cond: (sorted = 320211)
  Heap Fetches: 5
  Planning Time: 2.624 ms
  Execution Time: 1.343 ms
(5 rows)

hw5=# explain analyze select unsorted from table1 where unsorted=320211;
                                QUERY PLAN
-----
Index Only Scan using table1_unsorted_idx on table1 (cost=0.43..28.54 rows=6 width=4) (actual time=1.778..1.786 rows=4 loops=1)
  Index Cond: (unsorted = 320211)
  Heap Fetches: 4
  Planning Time: 0.113 ms
  Execution Time: 1.811 ms
(5 rows)
```

Q3.

```
explain analyze select sorted, rndm from table where sorted > 100 and sorted < 500 and rndm = 55;
explain analyze select unsorted, rndm from table where unsorted > 100 and unsorted < 500 and rndm = 55
```

```
hw5=# explain analyze SELECT sorted, rndm FROM table1 WHERE sorted > 100 and sorted < 500 AND rndm = 55;
                                QUERY PLAN
-----
Index Scan using index1 on table1 (cost=0.43..92.94 rows=1 width=8) (actual time=4.266..4.266 rows=0 loops=1)
  Index Cond: ((sorted > 100) AND (sorted < 500))
  Filter: (rndm = 55)
  Rows Removed by Filter: 1995
  Planning Time: 0.779 ms
  Execution Time: 4.288 ms
(6 rows)

hw5=# explain analyze SELECT unsorted, rndm FROM table1 WHERE unsorted > 100 and unsorted < 500 AND rndm = 55;
                                QUERY PLAN
-----
Bitmap Heap Scan on table1 (cost=44.37..7155.78 rows=1 width=8) (actual time=44.142..44.142 rows=0 loops=1)
  Recheck Cond: ((unsorted > 100) AND (unsorted < 500))
  Filter: (rndm = 55)
  Rows Removed by Filter: 2021
  Heap Blocks: exact=2005
  -> Bitmap Index Scan on index2 (cost=0.00..44.37 rows=1993 width=0) (actual time=3.830..3.830 rows=2021 loops=1)
        Index Cond: ((unsorted > 100) AND (unsorted < 500))
  Planning Time: 0.300 ms
  Execution Time: 44.224 ms
(9 rows)
```

clustered index가 4.288ms, non-clustered index가 44.244ms이다.

clustered index에서는 데이터가 sorted되어 있기 때문에 bitmap heap scan과 같은 disk access에서 시간을 생략하거나 단축할 수 있어서 더 효율적이다.

Q4.

```
explain analyze select sorted, rndm from table1 where sorted >1999231 and rndm=1005;
explain analyze SELECT sorted , rndm FROM table1 WHERE sorted < 1999231 AND rndm = 1005;
```

```
hw5=# explain analyze select sorted, rndm from table1 where sorted >1999231 and rndm=1005;
                                QUERY PLAN
-----
Index Scan using table1_sorted_idx on table1 (cost=0.43..153.84 rows=1 width=8) (actual time=0.653..0.653 rows=0 loops=1)
  Index Cond: (sorted > 1999231)
  Filter: (rndm = 1005)
  Rows Removed by Filter: 3840
  Planning Time: 0.139 ms
  Execution Time: 0.665 ms
(6 rows)
```

```
hw5=# explain analyze SELECT sorted , rndm FROM table1 WHERE sorted < 1999231 AND rndm = 1005;
                                QUERY PLAN
-----
Seq Scan on table1 (cost=0.00..253093.00 rows=100 width=8) (actual time=228.958..17110.581 rows=114 loops=1)
  Filter: ((sorted < 1999231) AND (rndm = 1005))
  Rows Removed by Filter: 9999886
  Planning Time: 0.149 ms
  JIT:
    Functions: 4
    Options: Inlining false, Optimization false, Expressions true, Deforming true
    Timing: Generation 3.073 ms, Inlining 0.000 ms, Optimization 0.827 ms, Emission 5.387 ms, Total 9.287 ms
  Execution Time: 17113.892 ms
(9 rows)
```

(1)는 0.665ms, (2)는 17113.892이다

우선 총 범위가 20000000인데, 1999231보다 큰 값을 조사하는 (1)이 1999231보다 작은 값을 조사하는 (2)보다 검사해야 할 row가 적기 때문에, (1)이 더 빠른 것이다.

B-tree and Hash

create two indexes

```
create index on table_btree using btree (recordid);
create index on table_hash using hash (recordid);
```

```
[hw5=# create index on table_btree using btree (recordid);
CREATE INDEX
[hw5=# create index on table_hash using hash (recordid);
CREATE INDEX
```

Q5.

```
explain analyze select * from tagble_btree where recordid=10001;
explain analyze select * from tagble_hash where recordid=10001;
```

```
[hw5=# explain analyze Select * from table_btree where recordid = 10001;
                                         QUERY PLAN
-----
Index Scan using table_btree_recordid_idx on table_btree  (cost=0.43..8.45 rows=1 width=49) (actual time=1.231..1.233 rows=1 loops=1)
  Index Cond: (recordid = 10001)
  Planning Time: 4.287 ms
  Execution Time: 1.261 ms
(4 rows)

[hw5=# explain analyze Select * from table_hash where recordid = 10001;
                                         QUERY PLAN
-----
Index Scan using table_hash_recordid_idx on table_hash  (cost=0.00..8.02 rows=1 width=49) (actual time=0.049..0.051 rows=1 loops=1)
  Index Cond: (recordid = 10001)
  Planning Time: 0.295 ms
  Execution Time: 0.070 ms
(4 rows)
```

모두 index scan을 사용하고, hash의 스캔이 더 빠른 것을 볼 수 있다.

Q6.

```
explain analyze select * from table_btree where recordid >250 and recordid <500;
explain analyze select * from table_noindex where recordid >250 and recordid <500;
```

```

[hw5=# explain analyze Select * from table_btree where recordid > 250 and recordid < 550;
]
QUERY PLAN
-----
Index Scan using table_btree_recordid_idx on table_btree (cost=0.43..15.78 rows=267 width=49) (actual time=0.452..1.58
4 rows=299 loops=1)
  Index Cond: ((recordid > 250) AND (recordid < 550))
  Planning Time: 3.819 ms
  Execution Time: 1.650 ms
(4 rows)

[hw5=# explain analyze Select * from table_hash where recordid > 250 and recordid < 550;
]
QUERY PLAN
-----
Seq Scan on table_hash (cost=0.00..253093.00 rows=1 width=49) (actual time=6.799..1497.806 rows=299 loops=1)
  Filter: ((recordid > 250) AND (recordid < 550))
  Rows Removed by Filter: 9999701
  Planning Time: 3.023 ms
  JIT:
    Functions: 2
    Options: Inlining false, Optimization false, Expressions true, Deforming true
    Timing: Generation 2.025 ms, Inlining 0.000 ms, Optimization 0.578 ms, Emission 3.374 ms, Total 5.977 ms
  Execution Time: 1500.018 ms
(9 rows)

```

hash에서 index scan이 아닌 seq scan을 사용하고, 더 많은 시간이 소비되는 것을 볼 수 있다.

Q7.

```

explain analyze update table_btree
set recordid=3
where recordid=2;

```

```

[hw5=# where recordid=2;
]
QUERY PLAN
-----
Update on table_btree (cost=0.43..8.45 rows=1 width=55) (actual time=1.363..1.363 rows=0 loops=1)
-> Index Scan using table_btree_recordid_idx on table_btree (cost=0.43..8.45 rows=1 width=55) (actual time=0.016..0
.017 rows=1 loops=1)
  Index Cond: (recordid = 2)
  Planning Time: 4.977 ms
  Execution Time: 2.874 ms
(5 rows)
...skipping...
]

```

```

explain analyze update table_noindex
set recordid=3
where recordid=2;

```

```

[hw5=# explain analyze update table_noindex
[hw5=# set recordid=3
[hw5=# where recordid=2;
]
]
]
QUERY PLAN
-----
Update on table_noindex (cost=0.00..228093.26 rows=1 width=55) (actual time=2172.172..2172.172 rows=0 loops=1)
-> Seq Scan on table_noindex (cost=0.00..228093.26 rows=1 width=55) (actual time=6.407..2171.457 rows=1 loops=1)
  Filter: (recordid = 2)
  Rows Removed by Filter: 9999999
  Planning Time: 1.119 ms
  JIT:
    Functions: 4
    Options: Inlining false, Optimization false, Expressions true, Deforming true
    Timing: Generation 2.020 ms, Inlining 0.000 ms, Optimization 0.348 ms, Emission 4.193 ms, Total 6.561 ms
  Execution Time: 2174.360 ms
(10 rows)

```

btree는 index scan로 1.119ms, noindex는 seq scan으로 2174.360ms 이다.

한 개의 key를 찾고 업데이트를 하는 것이니, btree가 더 효과적인 것을 확인할 수 있다.

Q8

```
explain analyze update table_btree
set recordidd =recordid *2
where recordid > 8000000;

explain analyze update table_noindex
set recordidd =recordid *2
where recordid > 8000000;
```

```
hw5=# explain analyze update table_btree
set recordid = recordid*2
where recordid>8000000;
                                QUERY PLAN
-----
Update on table_btree (cost=0.43..81708.38 rows=1979347 width=55) (actual time=23300.290..23300.290 rows=0 loops=1)
-> Index Scan using table_btree_recordid_idx on table_btree (cost=0.43..81708.38 rows=1979347 width=55) (actual time=1.215..2018.561 rows=1999999 loops=1)
    Index Cond: (recordid > 8000000)
    Planning Time: 0.094 ms
    Execution Time: 23300.320 ms
(5 rows)

hw5=# explain analyze update table_noindex
hw5=# set recordid=recordid*2
hw5=# where recordid>8000000;
                                QUERY PLAN
-----
Update on table_noindex (cost=0.00..233067.23 rows=1989587 width=55) (actual time=12134.881..12134.882 rows=0 loops=1)
-> Seq Scan on table_noindex (cost=0.00..233067.23 rows=1989587 width=55) (actual time=1108.439..1898.394 rows=1999999 loops=1)
    Filter: (recordid > 8000000)
    Rows Removed by Filter: 8000001
    Planning Time: 0.072 ms
    JIT:
        Functions: 4
        Options: Inlining false, Optimization false, Expressions true, Deforming true
        Timing: Generation 2.059 ms, Inlining 0.000 ms, Optimization 0.509 ms, Emission 4.388 ms, Total 6.956 ms
    Execution Time: 12137.024 ms
(10 rows)
```

btree 는 index scan으로 23300.320ms, noindex는 seq scan으로 12137.023ms이다
부분(한 개가 아닌 다수)을 업데이트를 하는 것은 noindex가 더 효율적인 것을 볼 수 있다.

Q9.

```
explain analyze update table_btree
set recordid = recordid*5;

explain analyze update table_noindex
set recordid = recordid*5;
```

```

[hw5=# explain analyze update table_btree
[hw5=# set recordid=recordid*5;
                                QUERY PLAN
-----
Update on table_btree (cost=0.00..248712.00 rows=10000000 width=55) (actual time=83121.647..83121.647 rows=0 loops=1)
  -> Seq Scan on table_btree (cost=0.00..248712.00 rows=10000000 width=55) (actual time=4.871..9542.640 rows=10000000
loops=1)
Planning Time: 3.420 ms
JIT:
  Functions: 2
  Options: Inlining false, Optimization false, Expressions true, Deforming true
  Timing: Generation 1.199 ms, Inlining 0.000 ms, Optimization 0.569 ms, Emission 3.281 ms, Total 5.049 ms
Execution Time: 83123.115 ms
(8 rows)

[hw5=# explain analyze update table_noindex
[hw5=# set recordid=recordid*5;
                                QUERY PLAN
-----
Update on table_noindex (cost=0.00..248712.00 rows=10000000 width=55) (actual time=34539.422..34539.422 rows=0 loops=1
)
  -> Seq Scan on table_noindex (cost=0.00..248712.00 rows=10000000 width=55) (actual time=5.701..12588.149 rows=10000
000 loops=1)
Planning Time: 2.786 ms
JIT:
  Functions: 2
  Options: Inlining false, Optimization false, Expressions true, Deforming true
  Timing: Generation 0.925 ms, Inlining 0.000 ms, Optimization 0.276 ms, Emission 3.254 ms, Total 4.455 ms
Execution Time: 34540.471 ms
(8 rows)

```

btree는 seq scan으로 83123.155ms, noindex는 seq scan으로 34540.471ms이다.
noindex의 효율이 더 많이 좋은 것을 확인할 수 있다.